Praneeth Papishetty

Storage engineer. Relational and NoSQL database systems.

Kirkland, WA | praneeth1962@gmail.com | (270) 438-6440 | linkedin.com/in/praneethpapishetty | praneethpapishetty.com

SUMMARY

I have worked with databases for 13 years. I started as an Oracle DBA. Then I ran Cassandra, ScyllaDB, Couchbase and MySQL NDB at AT&T. At Warner Bros. Discovery, I built the platform that deployed databases for HBO Max. Now I work on Vitess and MySQL at Slack.

EXPERIENCE

Slack

Mar 2026 – Present

Senior Software Engineer, Datastores Foundation | Remote

- I work on the platform layer under Slack's Vitess fleet: compute, packaging and configuration.
- I build tools that make Vitess and MySQL more reliable, and I automate routine maintenance.

Warner Bros. Discovery

Jan 2019 – Mar 2026

Sr. Staff Software Engineer, Database Reliability, HBO Max | Remote

- I built DIM, a self-service platform for database deployments. Provisioning time went from 3 weeks to 12 hours. Support requests went down by about 50%. Three other teams now use the same pattern.
- I co-designed the RDS major-version upgrade pipeline with DMS and CDC. Downtime was near zero. Manual work went down by about 80%.
- I wrote composite Terraform modules and a Checkov auto-review. Deployment time for EMEA and APAC went down by about 70%.
- I built a database connection tool with Homebrew, Docker and Steampipe. Support volume went down by about 30%.
- I redesigned how we deploy Aurora clusters and moved them to a multi-tenant model. AWS cost went down by about 80%.
- I automated Cassandra rebuilds and multi-datacenter scaling. Manual work went down by about 60%.
- I built a dashboard for HBO Max client metrics, for an internal project.

AT&T

Jan 2015 – Jan 2019

Database Analyst / DBA, eSUPPORT | Dallas, TX

- I fixed unpredictable compaction on DataStax Cassandra. TTLs ranged from seconds to days. The app team moved to predictable TTLs. I moved the tables to time-window compaction. Compaction time, CPU load and I/O load went down.
- I tuned and automated Cassandra repairs with Prometheus and Grafana. Repair time went down by about 80%.
- I put Oracle GoldenGate replication into production: one-way, two-way and cascade, with Veridata checks.
- I ran ScyllaDB, Couchbase, MariaDB and MySQL NDB clusters in production.
- I planned capacity for each iPhone launch, three months ahead.

Dollar Thrifty / Hertz

May 2013 – Jan 2015

Oracle DBA, Reservations | Tulsa, OK

- I ran more than 30 Oracle instances. One production database was 14 TB.
- I tuned SQL and indexes with AWR, ADDM and TKPROF. Peak-hour latency on bookings went down by about 50%.
- I archived old reservations from Oracle to MSSQL with SSIS.
- I led the schedule for an offshore team of four DBAs.

SKILLS

Datastores: Vitess, MySQL, PostgreSQL, Aurora, Oracle RAC, Cassandra, ScyllaDB, DynamoDB, Couchbase, ElastiCache, MongoDB, Redis
Storage internals: Compaction tuning, TTL design, Replication, Repairs, Query tuning, Indexing, Zero-downtime migrations
Infrastructure: AWS, GCP, Kubernetes, Terraform, Terragrunt, Argo Workflows, Jenkins, GitHub Actions
Languages: Python, Go, SQL, Shell, Groovy
Observability: Prometheus, Grafana, Splunk