

Praneeth Papishetty

Infrastructure & Reliability Engineer | Large-Scale Distributed Systems | Cloud Automation

Kirkland, WA | praneeth1962@gmail.com | (270) 438-6440 | linkedin.com/in/praneethpapishetty

Summary

Infrastructure engineer with 10+ years building and scaling large-scale cloud systems across AWS and GCP. Combines deep database expertise with modern infrastructure automation to deliver high-availability, secure, and cost-optimized platforms. Skilled at designing self-service systems, automating multi-region deployments, and leading complex migrations with minimal downtime. Adept at accelerating developer productivity through custom tools and AI-driven validation systems.

Key Achievements

- Cut global market deployment time by 70% via Terraform module refactoring and automation.
- Built Database Infra Manager (DIM), enabling fully automated database deployments and reducing support volume by 50%.
- Created DB connection automation tool for incident/ops use; improved response speed and cut related tickets by 30%.
- Delivered 80% AWS cost optimization via strategic architecture redesign.
- Automated multi-DC Cassandra rebuilds and scaling, reducing manual effort by 60%.

Core Skills

Infra & Distributed Systems: Kubernetes, AWS, GCP, Terraform, Argo Workflows, GitHub Actions, Ansible, Jenkins, IaC pipelines, HA/DR design, Cost Optimization

Data & Storage: PostgreSQL, Cassandra, DynamoDB, Couchbase, ScyllaDB, Redis, Presto, Zeppelin

Automation & Tooling: Python, Go, Shell, Groovy, Rules Engines, Self-Service Portals, CI/CD, Monitoring/Alerting (Prometheus, Grafana, Splunk)

Security & Governance: IAM, RBAC, audit logging, schema compliance, FinOps best practices

Collaboration: Tech leadership, stakeholder engagement, mentoring, incident response

Professional Experience

Warner Bros. Discovery — Sr. Staff Software Engineer (Infrastructure & Reliability) | Sep 2021 – Present

- Led automation of database and core infra deployments across AWS, integrating IaC, pipelines, and developer portals.
- Designed and built Database Infra Manager (DIM), a rules-driven, self-service deployment framework; accelerated provisioning cycles and reduced support requests ~50%.
- Built Python-based DB connection tool to abstract manual SSM bastion tunneling and login; leverages Steampipe to fetch cluster details and Dockerized psql/pgcli clients for direct connections. Reduced support volume by ~30% and improved incident response by eliminating manual endpoint and credential lookups.
- Engineered RDS EOL upgrade pipeline enabling near-zero downtime migrations with DMS & CDC; saved ~80% manual effort.
- Automated Cassandra cluster rebuilds and multi-DC scaling; reduced operational toil ~60%.

- Owned infra-deployment pipelines (Terraform, Terragrunt, Jenkins, Argo, GitHub Actions) across multi-region architectures.
- Built monitoring & alerting with Grafana, Prometheus, postgres-exporter for distributed systems.

HBO (via Metis Software Solutions) — Database Reliability Engineer | Jan 2019 – Sep 2021

- Developed Terraform/Terragrunt modules with automated testing, reducing AWS deployment time by ~50%.
- Created PR review automation for Terraform plans; cut deployment time by ~30%.
- Built CI/CD for database creation, migrations, and upgrades (Jenkins, GHA).
- Automated Grafana dashboards for multiple data stores; deployed postgres-exporter to Kubernetes.

AT&T; — Database Analyst/DBA | Jan 2015 – Jan 2019

- Maintained large-scale Oracle RAC, Cassandra (100+ nodes), Couchbase, ScyllaDB, PostgreSQL, and MySQL NDB Cluster environments.
- Automated Cassandra repairs (node exporter/Prometheus/Grafana), cutting repair time by ~80%.
- Implemented Oracle GoldenGate topologies (uni/bi-directional, cascade) and Veridata monitoring.
- Implemented Presto + Zeppelin for ad-hoc reporting over distributed stores; collaborated on data modeling.
- Managed MySQL NDB Cluster environments and implemented Oracle GoldenGate for heterogeneous replication.

Dollar Thrifty Automotive Group — Oracle DBA | May 2013 – Jan 2015

- Administered Oracle 10g/11g/RAC across prod and non-prod; supported 30+ instances including 14 TB prod DBs.
- Improved app performance ~50% through SQL and database tuning (AWR/ADDM/TKPROF/Explain Plan).
- Led offshore team of 4; automated backups/monitoring with shell scripting and DBA jobs.
- Executed Oracle→MSSQL data migration via SSIS; ensured integrity and minimal downtime.

Wilmington University — Intern, Electronic Medical Records Project | Dec 2012 – Apr 2013

- Built Oracle 11g schemas for EMR: tables, views, procedures, triggers; enforced RI via PK/FK/checks.
- Created ~30 tables and database objects; collaborated with UI team to integrate data layer.

Helms-Cravens Library (WKU) — Library Assistant | Sep 2011 – Dec 2011

- Managed catalog intake and inventory workflows; ensured accurate data entry and organization.

Education

Master's Degree in Information System Technology — Wilmington University (2013)

Certifications

Apache Cassandra Certified Professional